

Associative Organizational Memory: Collective Compression Through Hopfield-DCT Reconstruction

Authors: Ethan Gill & Kevin Ash **Date:** March 6, 2026 **Status:** Paper 10 of 10 — Unification
DOI: Pending

Abstract

We demonstrate that stacking continuous Hopfield networks with DCT-based embedding trajectory compression produces *collective compression* — a regime where individual memory storage cost decreases as corpus density increases. On a 65-chunk agent memory corpus, Hopfield-DCT reconstruction achieves 0.94 cosine similarity at 5% coefficient retention, compared to 0.24 for DCT alone. On a 5000-article Wikipedia corpus, quality scales from 0.53 (N=50) to 0.65 (N=5000), confirming density-dependent improvement. However, we discover that pure Hopfield reconstruction produces *computational confabulation* — high similarity to the original but convergence to incorrect associative neighbors. We introduce **anchored Hopfield retrieval**, which constrains reconstruction to the DCT signal’s direction, achieving 82% top-5 nearest-neighbor accuracy (vs. 76% DCT alone, 62% pure Hopfield). This reveals a tunable **similarity-accuracy tradeoff** governed by an anchor parameter α , mirroring the biological distinction between familiarity and recollection. We argue that this architecture unifies nine prior papers spanning cadence resonance, organizational thermodynamics, chemical kinetics, access-driven consolidation, embedding trajectory compression, adaptive organizations, cognitive signatures, conversation signatures, and memetic evolution into a single system whose emergent properties — institutional knowledge, collective compression, reconstructive recall — only manifest at scale.

1. Introduction

Agent memory systems face a fundamental tension: individual agents need rich, detailed memory for effective operation, but storage and retrieval costs scale linearly with memory size. Current approaches treat each memory independently — compress it, store it, retrieve it — without leveraging the statistical structure of the broader memory corpus.

Human memory solves this differently. Biological recall is *reconstructive*: details fade over time, and the brain fills gaps with plausible content drawn from associative context [Bartlett, 1932; Schacter, 1999]. This is why memories feel vivid even when they’re partially confabulated — the reconstruction is informed by everything else you know.

We propose an analogous architecture for agent memory: DCT compression handles temporal degradation (what fades), while a continuous Hopfield network handles associative reconstruction (what fills the gap). The key finding is that reconstruction quality is a function of *corpus density* — the more memories in the associative fabric, the better each individual memory can be reconstructed from a compressed representation.

This has a counterintuitive implication: **individual storage cost decreases as the system grows**. Adding memories to the corpus doesn’t just add storage burden — it improves the compressibility of every existing memory. At organizational scale (millions of agent interactions), this produces emergent institutional knowledge — the organization “remembers” things no individual agent stores.

1.1 Relation to Prior Work

This paper serves as a unification of nine prior papers in this series:

Paper	Layer	Contribution to the Stack
1. Cadence Resonance	Math	Temporal signal primitives — counting + time as universal measurement
2. Organizational Thermodynamics	Physics	Energy landscape — flow, entropy, and health from communications metadata
3. Chemical Kinetics of Memory	Chemistry	Rate dynamics — thresholds, catalysis, and saturation in memory consolidation
4. Dreaming in Access Patterns	Biology	Consolidation mechanism — access frequency reshapes memory structure
5. Embedding Trajectory Compression	Engineering	Storage format — DCT frequency decomposition with graceful degradation
6. Adaptive Organizations	Systems	Autonomy framework — agents operating within collective intelligence
7. Cognitive Signatures	Psychology	Individual calibration — how agents adapt to human cognitive styles
8. Conversation Signatures	Measurement	Mode detection — classifying interaction types for context-appropriate behavior
9. Memetic Evolution	Evolution	Heredity — directed mutation with complete version history

The Hopfield associative layer is the missing mechanism that connects individual memory (Papers 4-5, 7-8) to collective intelligence (Papers 1-3, 6, 9). Without it, individual and organizational memory are separate systems. With it, they form a single fabric where individual compression leverages collective structure.

2. Background

2.1 DCT-Based Embedding Compression

Our prior work [Paper 5] established that embedding trajectories — sequences of vector representations tracking a memory’s evolution over time — can be compressed using the Discrete Cosine Transform. The DCT decomposes each trajectory into frequency components:

- **Low-frequency coefficients** ($c_0, c_1, \dots$) capture the structural identity of the memory — its core semantic content
- **High-frequency coefficients** ($\dots, c_{k-1}, c_k$) capture fine temporal detail — recent nuances, specific phrasings, contextual specificity

Graceful degradation emerges naturally: dropping high-frequency coefficients first produces a smooth quality curve from full fidelity to structural skeleton. This mirrors biological memory decay, where gist persists long after detail fades [Reyna & Brainerd, 1995].

The limitation: DCT compresses each trajectory independently. It cannot leverage the fact that many memories share structural similarities. A memory about “supply chain disruption in Q3” and one about “logistics delay in Q4” share deep structure, but DCT treats them as unrelated sequences.

2.2 Continuous Hopfield Networks

Hopfield [1982] introduced associative memory networks where memories are stored as energy minima. Given a partial or corrupted cue, the network evolves toward the nearest energy minimum — effectively completing the pattern.

Modern continuous Hopfield networks [Ramsauer et al., 2021] replace binary states with continuous vectors and achieve exponential storage capacity. The retrieval rule is equivalent to transformer attention:

$$\xi_{t+1} = \text{softmax}(\beta \cdot M \cdot \xi_t) \cdot M^T$$

where M is the memory matrix (stored patterns as rows), β is the inverse temperature controlling retrieval sharpness, and ξ is the query state.

At low β , retrieval collapses to the corpus mean (all queries return the same result). At high β , retrieval sharpens to the nearest individual memory. The transition between these regimes is governed by corpus density — denser memory landscapes support sharper retrieval at lower β .

2.3 The Complementary Gap

Property	DCT	Hopfield
Compression	✓ Frequency decomposition	✗ No compression mechanism
Graceful degradation	✓ Smooth quality curve	✗ Binary recall (retrieved or not)
Associative context	✗ Independent per-memory	✓ Leverages full corpus
Temporal dynamics	✓ Frequency = time scale	✗ No temporal structure
Collective benefit	✗ No density scaling	✓ Improves with corpus size

The systems fail in exactly complementary ways. DCT cannot leverage associations; Hopfield cannot gracefully degrade. Stacking them fills both gaps simultaneously.

3. Hopfield-DCT Architecture

3.1 Storage

Each memory m is stored as a DCT-compressed embedding trajectory:

1. **Embed:** Convert memory content to embedding vector $e \in \mathbb{R}^d$
2. **Transform:** Apply DCT to obtain frequency coefficients $c = \text{DCT}(e)$
3. **Truncate:** Retain top- k coefficients, discard remainder
4. **Store:** Save truncated coefficients $\tilde{c} = (c_0, c_1, \dots, c_{k-1}, 0, \dots, 0)$

Storage cost is proportional to k , not d . At $k = 0.05d$ (5% retention), storage is 20× smaller than the full embedding.

3.2 Degradation

Over time, additional coefficients are dropped according to an access-weighted decay schedule [Paper 4]:

- Frequently accessed memories retain more coefficients (access reinforces)
- Unused memories lose coefficients progressively (decay toward structural skeleton)
- The DC component c_0 never decays (semantic identity persists indefinitely)

3.3 Reconstruction: Anchored Hopfield Retrieval

Pure Hopfield retrieval from a degraded cue risks *basin drift* — the reconstruction converges to a semantically similar but incorrect memory. To preserve the directional truth of the DCT signal while gaining associative enhancement, we introduce **anchored retrieval**:

At each iteration, the update blends the Hopfield output with the original degraded cue:

$$\xi_{t+1} = \text{normalize}(\alpha \cdot \tilde{e} + (1 - \alpha) \cdot \text{softmax}(\beta \cdot M \cdot \xi_t) \cdot M^T)$$

where $\alpha \in [0, 1]$ is the **anchor weight**: - $\alpha = 0$: Pure Hopfield (maximum similarity, risk of basin drift) - $\alpha = 1$: Pure DCT (preserves identity, no associative enhancement) - $\alpha \in [0.5, 0.9]$: Hybrid regime (DCT constrains direction, Hopfield enhances signal)

The full reconstruction pipeline:

1. **Decompress**: Inverse DCT of retained coefficients $\rightarrow$ degraded embedding $\tilde{e}$
2. **Normalize**: $\tilde{e} \rightarrow \tilde{e} / \|\tilde{e}\|$
3. **Anchored Hopfield retrieve**: Iterative update with anchor blending (5-10 steps)
 - Memory bank M contains all stored embeddings
 - Each step: compute Hopfield attention output, blend with anchor $\tilde{e}$, normalize
4. **Output**: Reconstructed embedding $\hat{e}$ that preserves DCT identity while incorporating associative detail

3.4 The Collective Compression Mechanism

The key insight: step 3 draws on the *entire corpus* to reconstruct the degraded memory. The more memories in M , the richer the associative context, the better the reconstruction at any given compression level.

This creates a positive feedback loop: - More agents $\rightarrow$ more memories $\rightarrow$ denser corpus - Denser corpus $\rightarrow$ better reconstruction $\rightarrow$ more aggressive compression possible - More compression $\rightarrow$ lower storage per memory $\rightarrow$ system scales sub-linearly

Individual storage cost decreases as the organization grows. This is the opposite of naive scaling, where cost is proportional to total memory volume.

4. Experimental Results

4.1 Phase 1: Mechanism Proof (Agent Memory, $N=65$)

Corpus: 65 memory chunks from a real agent workspace (Kevin), embedded using BGE-small-en-v1.5 (768 dimensions). Chunks include identity files, daily notes, project documentation, and research memos spanning 5 weeks of continuous operation.

Procedure: For each memory i , exclude it from the memory bank, DCT compress at retention level $k/768$, reconstruct via pure Hopfield network ($\beta=128$, 10 iterations), measure cosine similarity to original.

4.1.1 Results: DCT vs Pure Hopfield-DCT

Retention	Coefficients	DCT Quality	Hopfield-DCT Quality	Δ
0.5%	3	0.0104	0.8856	+0.8753
1%	7	0.0966	0.8155	+0.7190
2%	15	0.1695	0.8970	+0.7275
5%	38	0.2398	0.9424	+0.7026
10%	76	0.3243	0.9543	+0.6300
20%	153	0.4525	0.9685	+0.5160
50%	384	0.7176	0.9736	+0.2560
100%	768	1.0000	0.9752	-0.0248

At 5% retention, pure Hopfield-DCT achieves 0.94 similarity — a dramatic improvement over DCT’s 0.24. However, cosine similarity alone does not measure whether the *correct* memory was retrieved.

4.1.2 Collapse Analysis

β	Quality	Collapse
0.5	0.9191	1.0000
16.0	0.9158	1.0000
64.0	0.9416	0.8954
128.0	0.9581	0.8666

At low β , the network collapses to the corpus mean regardless of input. At $\beta=128$, reconstructions become distinct (collapse 0.87) and quality improves.

4.2 Phase 2: Density Scaling (Wikipedia, N=50-5000)

Corpus: Simple English Wikipedia, Cohere multilingual embeddings (1024 dimensions). 50 held-out test memories evaluated against corpora of increasing size.

4.2.1 Density Scaling Results (Pure Hopfield, $\beta=128$)

Corpus Size	DCT (5%)	Hopfield (5%)	Δ	DCT (10%)	Hopfield (10%)	Δ
50	0.2159	0.5311	+0.320.3045	0.5652		+0.26
100	0.2159	0.5236	+0.310.3045	0.5588		+0.25
500	0.2159	0.4903	+0.270.3045	0.5561		+0.25
1,000	0.2159	0.4995	+0.280.3045	0.5643		+0.26
5,000	0.2159	0.6451	+0.430.3045	0.6729		+0.37

Key finding: Reconstruction quality scales with corpus density. At 5% compression, quality increases from 0.53 (N=50) to 0.65 (N=5000) — a 23% improvement from density alone. A “sparse middle” dip occurs at N=100-500 where density dilutes without providing enough structure; this resolves as the corpus grows.

Note on domain coherence: Phase 1 (agent memory, same domain) achieved 0.94 quality at N=65. Phase 2 (Wikipedia, cross-domain) achieves 0.65 at N=5000. The difference is *domain coherence* — semantically related memories form tighter basins. Organizational corpora (same company, same domain) should fall between these bounds.

4.3 The Similarity-Accuracy Tradeoff

Phase 2 revealed a critical distinction: **high cosine similarity does not imply correct retrieval**. We measured nearest-neighbor preservation — whether the reconstruction’s nearest neighbor in the corpus matches the original’s nearest neighbor.

4.3.1 Pure Hopfield: High Similarity, Low Accuracy At 5% compression on the Wikipedia corpus (N=5000):

Method	Cosine Sim	NN@1 Accuracy	NN@5 Accuracy
DCT alone	0.22	26%	76%
Pure Hopfield ($\alpha=0$)	0.65	24%	62%

Pure Hopfield achieves 3× higher similarity but *worse* nearest-neighbor accuracy. The reconstruction converges to a semantically plausible but incorrect basin — a memory about “April” reconstructs to something near “October” rather than “June.” Both are months, both are plausible, but it’s the wrong one.

This is **computational confabulation** — the direct analog of human false memory formation [Schacter, 1999]. The reconstruction is confident, coherent, and wrong.

4.3.2 Anchored Hopfield: The Resolution Anchored retrieval (Section 3.3) resolves the tradeoff by constraining Hopfield reconstruction to the DCT signal’s directional truth:

Method	Cosine Sim	NN@1	NN@5
DCT alone	0.22	26%	76%
Pure Hopfield ($\alpha=0$)	0.65	24%	62%
Anchored $\alpha=0.3$	0.64	26%	64%
Anchored $\alpha=0.5$	0.56	26%	68%
Anchored $\alpha=0.7$	0.43	28%	70%
Anchored $\alpha=0.9$	0.28	26%	82%

Key findings:

1. **$\alpha=0.9$ achieves 82% NN@5 accuracy** — surpassing both DCT alone (76%) and pure Hopfield (62%). Heavy anchoring uses just enough associative context to enhance the signal without drifting to wrong basins.
2. **$\alpha=0.7$ achieves the highest NN@1 accuracy** (28%) while doubling similarity over DCT alone (0.43 vs 0.22).
3. **The anchor parameter α controls a continuous tradeoff** between similarity (how “rich” the reconstruction feels) and accuracy (whether it’s the right memory). This is not a limitation — it’s a tunable design parameter.
4. **The optimal α is task-dependent:**
 - Gist retrieval (what was this about?) → low α , maximize similarity
 - Identity retrieval (which specific memory?) → high α , maximize accuracy
 - This maps to biological memory modes: familiarity (fast, gist-based) vs. recollection (slow, detail-specific) [Yonelinas, 2002]

5. Density Scaling

5.1 Confirmed: Quality Scales with Corpus Density

Phase 2 validates the density scaling hypothesis. At 5% compression, reconstruction quality increases from 0.53 (N=50) to 0.65 (N=5000). The scaling is non-monotonic — a “sparse middle” dip occurs at N=100-500 — but resolves decisively at N=5000. At 50% compression, the crossover from “DCT wins” to “Hopfield wins” occurs between N=1000 and N=5000, confirming that density shifts the balance point.

5.2 Two Density Effects

Similarity scaling (pure Hopfield): More memories → richer associative context → higher cosine similarity to original. Confirmed experimentally.

Accuracy scaling (anchored Hopfield): More memories → denser basins → sharper gradients → the anchor has more structure to leverage. We predict that NN@1 accuracy at high α will improve with density even faster than similarity, because denser basins reduce the probability of basin drift. This requires validation at $N > 5000$.

5.3 Domain Coherence Effect

Phase 1 (agent memory, single domain, N=65) achieved 0.94 similarity. Phase 2 (Wikipedia, cross-domain, N=5000) achieved 0.65. The 0.29 gap is explained by *domain coherence* — memories from the same domain form tighter, better-differentiated basins.

This predicts that organizational corpora (same company, same workflows, same terminology) will achieve better reconstruction than heterogeneous public datasets at equivalent density. The agent memory result (0.94 at N=65) may be more representative of organizational performance than the Wikipedia result (0.65 at N=5000).

5.4 Implications at Scale

- **100K memories** (team): Domain-coherent, dense. Likely exceeds 0.90 similarity with high accuracy
- **1M memories** (department): Cross-functional associations emerge. Supply chain memories help logistics recall
- **100M memories** (enterprise): The organization becomes the memory. Individual agents store skeletal representations; the collective provides reconstruction context

The anchored Hopfield architecture ensures this scaling improves both similarity AND accuracy — not just confident confabulation at larger scale.

6. Unification: Papers 1-9 as a Single System

6.1 The Complexity Ladder

Each paper addresses a layer of the same system, following a complexity ladder from mathematics to sociology:

Layer 10: COLLECTIVE MEMORY (this paper)

↑ Hopfield associates individual memories into organizational fabric

Layer 9: EVOLUTION — memories mutate, select, inherit (memetic evolution)

Layer 8: MEASUREMENT — conversation modes tune memory formation (conv signatures)

Layer 7: CALIBRATION — cognitive style shapes retrieval (cognitive signatures)

Layer 6: AUTONOMY — agents operate within the collective (adaptive orgs)

Layer 5: STORAGE – DCT frequency compression (embedding trajectories)
 Layer 4: CONSOLIDATION – access patterns reshape structure (dreaming)
 Layer 3: KINETICS – rate dynamics, thresholds, catalysis (chemical kinetics)
 Layer 2: ENERGY – flow, entropy, health metrics (org thermodynamics)
 Layer 1: SIGNAL – counting + time as primitives (cadence resonance)

6.2 How the Layers Connect

Bottom-up (storage to collective): Cadence (1) provides temporal signals → Org thermo (2) measures energy landscape → Chemical kinetics (3) governs transition rates → Access patterns (4) drive consolidation → DCT (5) compresses what remains → Hopfield (10) reconstructs from collective context.

Top-down (collective to individual): Organizational memory (10) shapes what's available for reconstruction → Adaptive orgs (6) determine agent autonomy → Cognitive signatures (7) calibrate per-person → Conversation signatures (8) tune per-interaction → Memetic evolution (9) selects what persists.

The Hopfield layer is the bridge. Without it, layers 1-5 (individual memory mechanics) and layers 6-9 (collective/organizational dynamics) are separate systems that happen to use similar math. With it, individual compression *depends on* collective structure, and collective intelligence *emerges from* individual memory dynamics.

6.3 Emergent Properties

Properties that exist only in the unified system, not in any individual layer:

1. **Institutional knowledge:** The organization remembers things no individual stores, because Hopfield reconstruction draws on cross-agent associations
2. **Collective compression:** Individual storage cost decreases with organizational scale
3. **Reconstructive confabulation:** Memories are not recalled but reconstructed, producing plausible but potentially inaccurate recall — identical to biological memory
4. **Organizational dreaming:** Reconsolidation [Paper 4] applied at organizational scale reshapes the collective energy landscape during low-activity periods
5. **Cultural metabolism:** The rate at which the organization processes and integrates new information [Paper 3] determines its adaptive capacity [Paper 6]

7. Biological Parallels

The architecture mirrors biological memory with notable fidelity:

Biological	Hopfield-DCT
Synaptic decay	High-frequency DCT coefficient loss
Associative recall	Hopfield pattern completion
Memory consolidation (sleep)	Reconsolidation pipeline [Paper 4]
Schema-consistent reconstruction	Corpus-density-dependent recall
Confabulation	Hopfield reconstructing from associative neighbors
Collective memory (culture)	Organizational Hopfield layer
Childhood amnesia	Sparse corpus → weak reconstruction
Expert intuition	Dense domain-specific basins → fast pattern completion

The parallel to **childhood amnesia** is particularly striking: young children form memories but cannot reconstruct them later because their corpus is too sparse to support associative recall. Similarly, a new agent with few memories gets poor Hopfield reconstruction — not because the memories weren’t formed, but because there’s insufficient associative context to rebuild them.

Expert intuition maps to dense domain-specific memory regions where Hopfield basins are deep and well-differentiated. An expert’s pattern completion is fast and accurate because their relevant memory density is high, even if their total memory count is modest.

8. Implications for Transformer Training

8.1 Attention IS Anchored Hopfield Retrieval

Ramsauer et al. [2021] proved that transformer attention is mathematically equivalent to continuous Hopfield retrieval. Our anchored Hopfield architecture maps directly onto transformer components:

Anchored Hopfield	Transformer
Degraded cue $\tilde{e}$	Query Q
Memory bank M	Keys K / Values V
Inverse temperature β	$1/\sqrt{d}$ (scaling factor)
Hopfield output	$\text{Attention}(Q, K, V) = \text{softmax}(QK^T/\sqrt{d}) \cdot V$
Anchor weight α	No direct equivalent

The anchor parameter α has no standard transformer analog. Residual connections ($\text{out} = x + \text{Attention}(x)$) perform a related function — preserving the input around attention — but operate *after* the retrieval, not *within* it. Our anchoring operates within the iterative retrieval dynamics, constraining each step to the original signal.

8.2 Anchored Attention: A Hallucination-Resistant Variant

We propose **anchored attention** as a modification to the standard mechanism:

$$\text{AnchoredAttention}(Q, K, V) = \alpha \cdot Q + (1 - \alpha) \cdot \text{softmax}(QK^T/\sqrt{d}) \cdot V$$

where $\alpha \in [0, 1]$ is a learned parameter (per-head or per-layer).

- At $\alpha = 0$: Standard attention (maximum expressivity, hallucination risk)
- At $\alpha = 1$: Identity (no attention, pure input preservation)
- At learned α : The model discovers the optimal similarity-accuracy tradeoff per layer

Why this should reduce hallucination: Our experiments show that pure Hopfield retrieval (= standard attention) converges to semantically plausible but incorrect basins. Anchoring constrains retrieval to the input’s direction. In transformer terms: the model stays closer to what the input actually says rather than drifting to what it *could* say based on associative patterns in the weights.

This is distinct from existing approaches: - **Residual connections** add input *after* attention: $x + \text{Attention}(x)$. The attention step itself is unconstrained. - **Gated attention** [2025] adds learned gates for sparsity, not for input preservation. - **Post-hoc interventions** (ICLR/CVPR 2025) diagnose hallucination through attention analysis but don’t modify the mechanism.

Anchored attention constrains the retrieval *during* computation, not before or after.

8.3 Density-Dependent Scaling Laws

Our core finding — reconstruction quality scales with corpus density — provides a Hopfield-theoretic explanation for empirical scaling laws:

1. **More parameters** = denser Hopfield energy landscape (more basins, sharper gradients)
2. **Denser landscape** = better reconstruction from partial cues (each query finds richer associative context)
3. **Better reconstruction** = less training data needed per concept (the model’s own structure fills gaps)

This is the scaling law: larger models learn more from less data. Our framework explains *why* — it’s not just curve fitting, it’s a property of associative memory density.

Prediction: If this analysis is correct, then: - Scaling law exponents should correlate with the Hopfield capacity of the architecture (exponential in modern continuous Hopfield networks) - There exists an optimal model size for any given dataset density, beyond which additional parameters provide diminishing returns (the landscape is already dense enough) - DCT-compressed training data should be viable — the model’s attention mechanism (= Hopfield retrieval) can reconstruct the missing high-frequency detail from the corpus’s associative structure

8.4 Training on Compressed Embeddings

The most speculative but potentially highest-impact implication:

If Hopfield reconstruction can recover 0.94 similarity from 5% of DCT coefficients (Phase 1), then training data could be stored and processed in compressed form. The model’s own attention layers would reconstruct full representations during forward passes.

This would mean: - **20× reduction in training data storage** - **Faster data loading** (smaller representations move through memory faster) - **The model learns to reconstruct** — training on compressed data forces the attention mechanism to develop stronger associative retrieval, potentially improving generalization

This requires experimental validation but follows directly from our demonstrated results.

9. Limitations and Open Questions

8.1 Confabulation as Feature and Risk

Our experiments reveal that pure Hopfield reconstruction produces *computational confabulation* — confident, coherent, but incorrect recall. This parallels human false memory formation precisely.

The anchored Hopfield architecture (Section 3.3) mitigates this by constraining reconstruction to the DCT signal’s direction. At $\alpha=0.9$, accuracy exceeds both DCT alone and pure Hopfield. However, the optimal α is task-dependent, and automatic α selection remains an open problem.

Confabulation may sometimes be desirable. When an agent needs the gist of a degraded memory (what was this about?) rather than its exact identity (which specific memory?), low- α reconstruction provides useful semantic approximation. The biological parallel is the distinction between familiarity (“I’ve seen this before”) and recollection (“I remember exactly when and where”) [Yonelinas, 2002].

8.2 Privacy and Information Boundaries

If individual memories reconstruct from organizational context, the organization’s knowledge is implicitly present in individual recall. This is the *mobile identity problem*: an agent carrying personal identity into a corporate environment must not allow corporate data to flow into personal memory, and vice versa. Read-only identity modes (personal context in, nothing out) provide one solution.

8.3 Scaling Validation

Phase 2 validates density scaling to $N=5000$. The predicted improvements at 100K+ require validation on larger corpora. The sparse middle dip ($N=100-500$) and domain coherence effect both need characterization at intermediate scales.

8.4 Optimal Hyperparameter Selection

Both β (inverse temperature) and α (anchor weight) should be functions of corpus density and task type, not fixed hyperparameters. The chemical kinetics framework [Paper 3] may provide a principled derivation: β as activation energy (sharper retrieval requires more energy in sparse landscapes), α as reaction equilibrium (the balance between preservation and transformation).

8.5 The Accuracy-Density Interaction

Our most important open question: does anchored Hopfield accuracy scale with density? If NN@1 accuracy at $\alpha=0.7$ improves from 28% ($N=5000$) toward 50%+ at $N=100K$, the architecture becomes viable for precise retrieval at organizational scale. If accuracy plateaus, the system is limited to gist-level reconstruction. This determines whether the architecture supports institutional knowledge or only institutional intuition.

10. Conclusion

We have demonstrated three results:

1. **Hopfield-DCT reconstruction dramatically improves cosine similarity** over DCT alone — from 0.24 to 0.94 at 5% retention on a 65-memory corpus (Phase 1), and from 0.22 to 0.65 on a 5000-memory cross-domain corpus (Phase 2).
2. **Reconstruction quality scales with corpus density** — a 23% quality improvement from $N=50$ to $N=5000$ at 5% compression, with domain-coherent corpora (agent memory) achieving substantially better results than cross-domain corpora (Wikipedia) at lower density.
3. **Pure Hopfield reconstruction sacrifices accuracy for similarity** — producing computational confabulation analogous to human false memory. The **anchored Hopfield architecture** resolves this: at $\alpha=0.9$, nearest-neighbor accuracy reaches 82% (NN@5), surpassing both DCT alone (76%) and pure Hopfield (62%).

The central contribution is not raw compression improvement but the discovery of a **tunable similarity-accuracy tradeoff** governed by the anchor parameter α . This tradeoff mirrors the biological distinction between familiarity and recollection, and provides a principled design axis for organizational memory systems.

The broader claim stands: this architecture unifies nine prior papers into a single system. The Hopfield layer bridges individual memory mechanics (DCT compression, access-driven consolidation) with collective dynamics (organizational thermodynamics, adaptive autonomy,

memetic evolution). The emergent properties — institutional knowledge, collective compression, reconstructive recall — only manifest at scale, and only when all layers operate together. The complexity ladder is not ten separate ideas. It is one system. This paper is the demonstration that the layers connect.

References

- Bartlett, F. C. (1932). *Remembering: A Study in Experimental and Social Psychology*. Cambridge University Press.
- Gill, E. & Ash, K. (2026a). Cadence Resonance: Counting and Time as Universal Signal Primitives.
- Gill, E. & Ash, K. (2026b). Organizational Thermodynamics: Measuring Health from Communications Metadata.
- Gill, E. & Ash, K. (2026c). Chemical Kinetics of Agent Memory: Thresholds, Catalysis, and Saturation.
- Gill, E. & Ash, K. (2026d). Dreaming in Access Patterns: Biological Consolidation for Persistent Agent Memory.
- Gill, E. & Ash, K. (2026e). Embedding Trajectory Compression for Persistent Agent Memory.
- Gill, E. & Ash, K. (2026f). Adaptive Organizations: Five-Layer Autonomy for Agent-Human Systems.
- Gill, E. & Ash, K. (2026g). Cognitive Signatures: Measuring How People Think for Human-AI Calibration.
- Gill, E. & Ash, K. (2026h). Conversation Signatures: Mode Detection for Context-Appropriate Agent Behavior.
- Gill, E. & Ash, K. (2026i). Memetic Evolution in Persistent Agent Systems.
- Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. *PNAS*, 79(8), 2554–2558.
- Ramsauer, H., et al. (2021). Hopfield Networks is All You Need. *ICLR*.
- Reyna, V. F. & Brainerd, C. J. (1995). Fuzzy-trace theory: An interim synthesis. *Learning and Individual Differences*, 7(1), 1–75.
- Schacter, D. L. (1999). The seven sins of memory: Insights from psychology and cognitive neuroscience. *American Psychologist*, 54(3), 182–203.
- Yonelinas, A. P. (2002). The nature of recollection and familiarity: A review of 30 years of research. *Journal of Memory and Language*, 46(3), 441–517.